

XNDL


Optimizers & regularization

Dario Garcia Gasulla
dario.garcia@bsc.es

Inertia in Optimizers

Momentum:


- ❖ Add a fraction of the previous gradient. *Inertia*.
- ❖ Decaying weight parameter. *Friction*.
- ❖ Faster, smoother convergence



Nesterov:

- ❖ Gradient computed after inertia
- ❖ See the slope ahead
- ❖ Faster convergence

momentum
gradient
update



Adaptative LR Optimizers

Adagrad:

- ❖ Apply LR to parameter-wise gradients (adaptative)
- ❖ Considering all past updates
- ❖ High LR for infrequent ones. Low LR for frequent ones.
- ❖ Good for sparse data.

Issues:

- ❖ Requires initial global LR
- ❖ Vanishing LR (monotonically decreasing)

Adaptative LR Optimizers

Adadelta:

- ❖ Use effective LR instead (past param. update / current gradient)
- ❖ Set a max. window, implemented as a decay avg.
- ❖ Requires decay rate (0.9?)

Adam:

- ❖ Momentum (Decaying avg of past gradients, mean, beta1)
- ❖ Adadelta (Decaying avg of past squared gradients, variance, beta2)

Nadam:

- ❖ Adadelta + Nesterov

AMSGrad, AdaMax, AdamW, ...

Regularization

Norm based (L1, L2):

- ❖ Add penalties to the value of parameters or activations
- ❖ Layer-specific
- ❖ Added to loss

```
from tensorflow.keras import regularizers

layer = layers.Dense(
    units=64,
    kernel_regularizer=regularizers.L1L2(l1=1e-5, l2=1e-4),
    bias_regularizer=regularizers.L2(1e-4),
    activity_regularizer=regularizers.L2(1e-5)
)
```

Optimizer Specific:

- ❖ Weight decay
- ❖ Gradient clipping (clipnorm, clipvalue)

```
tf.keras.optimizers.SGD(
    learning_rate=0.01,
    momentum=0.0,
    nesterov=False,
    weight_decay=None,
    clipnorm=None,
    clipvalue=None,
```

Regularization

- ❖ Mean of activations approximates 0, deviation of activations apprx. 1
- ❖ Contain trainable params (depends on distribution!)

LayerNorm

- ❖ Transform activations sample-wise

```
tf.keras.layers.LayerNormalization(
```

BatchNorm

- ❖ Transform activations batch-wise.

```
tf.keras.layers.BatchNormalization(
```

Testing the environment

1. Download Fashion MNIST dataset:

<https://github.com/zalandoresearch/fashion-mnist>

2. Upload it to the cluster

```
scp your_local_fashion_mnist.gz nct...@dt01.bsc.es
```

3. Place them here: ~/.keras/datasets

create directories first!

Files

- ❖ `lab7_code.py`
example of code for training a MLP on fashion dataset
- ❖ `launcher.sh`
example of script for submitting a job to execute *code.py*
- ❖ Submit job:
`sbatch launcher.sh`

Tasks

- ❖ Test different optimizers & hyperparameters
 - Convergence speed
- ❖ Try different normalization methods
 - Effect on overfitting
 - Convergence speed

Steps

- ❖ Test different optimizers & hyperparameters
 - Convergence speed
- ❖ Try different normalization methods
 - Effect on overfitting
 - Convergence speed

Baby steps

- ❖ Data nature, dimensionality, loss function, output layer + baseline
- ❖ Fix a batch size (stochasticity vs efficiency)
- ❖ Tune early stop, LR, activation funct., momentum while...
 - Start small **& underfit**
 - Grow **& overfit**
 - Regularize **& fit**

Dario Garcia-Gasulla (BSC)

dario.garcia@bsc.es

